Polisetti Venkata Ganga Hemanth Kumar

+91-7013341198 hemanthpolisetti7@gmail.com hemanth-kumar-polisetti HemanthPolisetti
Portfolio
Hyderabad, Telangana, India

Professional Summary

ETL Developer with 2+ years of production experience building integration codebases that connect payroll, HR, and business systems as both source and destination across REST, SOAP, webhooks, GraphQL, and OAuth2, processing 500K+ records daily in JSON, XML, Parquet, CSV, and Excel.
Reduced data processing time by 40% via CDC patterns and improved data accuracy by 35% through automated data-quality frameworks, with deployments shipped through CI/CD and Azure Key Vault.
Deliberately moving into **Azure Data Engineering**, with hands-on experience in Databricks and PySpark.

Technical Skills

Programming	PySpark, Python, SQL
Cloud	Azure Data Lake, Azure Blob Storage, Azure Data Factory, Azure Synapse, Azure Key Vault, Databricks
Databases	SQL Server, PostgreSQL, Redis
Integration	REST APIs, SOAP APIs, Webhooks, GraphQL, OAuth2
Data Formats	JSON, XML, Parquet, CSV, Excel
Tools	Docker, Git, GitHub, Apache Airflow, Apache Spark, Apache Kafka, CI/CD Pipelines, Claude Code
Concepts	ETL/ELT, CDC, Incremental Loading, Data Quality, Data Transformation (Mapping, Schema Reconciliation, Aggregation), Agile, Event-Driven Architecture

Professional Experience

ETL Developer **Feb 2024 – Present**

BenditoSoftTech *Hyderabad, India*

- Architected end-to-end ETL pipelines using PySpark, SQL, and Azure Data Factory, integrating REST, SOAP, and webhook-based APIs as both source and destination systems, processing 500K+ records daily into Azure Data Lake.
- Implemented incremental loading using Change Data Capture (CDC) patterns, reducing data processing time by 40% and optimizing cloud storage costs.
- Applied data transformation techniques including field mapping, format conversion, schema reconciliation, and aggregation across JSON, XML, Parquet, CSV, and Excel payloads.
- Developed automated data-quality frameworks covering schema compliance, null handling, and duplicate detection, improving data accuracy by 35%.
- Implemented bidirectional financial data synchronization involving job codes and journal entries between ERP and payroll systems, and built replay-safe batch-queue patterns for reliable recovery from partial failures.
- Maintained 99.5% pipeline uptime through proactive monitoring, performance tuning, and incident resolution across 15+ production pipelines.
- Shipped changes through CI/CD pipelines with Azure Key Vault-managed secrets, containerized workloads with Docker across dev, staging, and production environments, and promoted releases through structured sandbox-to-production testing; also supported OCR/RPA automation workflows.
- Collaborated with Data Analysts, BI teams, and Product Managers in an Agile environment to deliver business-aligned data solutions.

Key Projects

Enterprise Payroll & HR Integration Platform *Production*
PySpark, Azure Data Lake, Azure Data Factory, SQL Server, REST, SOAP, GraphQL, OAuth2, CI/CD, Azure Key Vault

- Designed and maintained a large integration codebase ingesting data from 5+ REST/SOAP APIs and 3 SQL databases into Azure Data Lake, processing 2M+ records weekly across systems acting as both source and destination.
- Traced and fixed a staging bug that reduced an intermediate dataset's memory footprint by roughly 40%; implemented data transformations including cleansing, normalization, and aggregation across JSON, XML, and CSV payloads with automated alerts for anomalies and schema drift.

- Deployed application changes through CI/CD pipelines with Azure Key Vault-backed secrets management.

Real-Time Stock Market Streaming Pipeline *In Progress*
Kafka, PySpark Structured Streaming, Delta Lake, ADLS Gen2, Docker

- Consumed live stock data from the Yahoo Finance API via Kafka, processed it with PySpark Structured Streaming, and computed rolling VWAP and volatility per ticker over 30-second micro-batch windows.
- Used Kafka checkpointing for exactly-once semantics, with Delta Lake output on ADLS Gen2.
- Fully containerized Kafka, Zookeeper, and PySpark using Docker Compose.

Reusable API Data Ingestion Framework *Complete*
PySpark, Python, REST APIs, Azure, Airflow

- Developed a modular framework for REST API ingestion with pagination, retry logic, and error handling, reducing onboarding time by 60%.
- Created a parameterized configuration system enabling rapid onboarding of new data sources without code changes.
- Orchestrated workflows with Apache Airflow; implemented a pytest suite with 85%+ coverage and GitHub Actions CI on every push.

Education

Bachelor of Technology in Mechanical Engineering **2019 – 2023**

Rajiv Gandhi University of Knowledge Technologies, Ongole, India *86%*

Certifications

- DevOps Essentials – CI/CD over AWS Cloud
- Complete Python Developer – Zero to Mastery